

HW3 Q7,8

Muhammad Ahmed Chaudhry

5/31/2021

Question 7

a) Here we code the function to implement SGD with mini-batches:

```
stochastic_gd <- function(X,y,epsilon,batch_size,epochs){  
  #' Function written as per required inputs and outputs  
  #' @param X: matrix, data-set of input variable values of size N*n  
  #' @param y: vector of responses for each given set of input variable values, size N*1  
  #' @param epsilon: numeric, value for the learning rate at each step  
  #' @param B: batch size, number of observations to use at each step when computing gradient  
  #' @param epochs: integer (positive), number of passes through the entire data-set  
  #' @return squared_loss: matrix, of squared error loss values, one for each step, size epochs*N/B  
  
  B<-batch_size  
  #print(B)  
  N<-dim(X)[1]  
  #print(N)  
  n<-dim(X)[2]  
  #print(n)  
  #indices<-seq(1,N/B,length.out=N)  
  #print(indices)  
  beta<-matrix(0,nrow=n,ncol=1)  
  squared_loss<-matrix(NA,nrow=epochs,ncol=N/B)  
  # loss_0<-1/N*t(y)%*%y  
  # squared_loss[1,]<-rep(loss_0,N/B)  
  
  for (j in 1:epochs){  
    loss<-c(rep(0,N/B))  
    for (i in 1:(N/B)){  
      start<-(i-1)*B+1  
      #print(start)  
      end<-start+B-1  
      #print(end)  
      predictors<-as.matrix(X[start:end,])  
      response<-y[start:end]  
      if (B==1){  
        gradient<-2/N*(as.matrix(predictors)%*(response-t(as.matrix(predictors))%*beta))  
        #print(gradient)  
        #print(dim(gradient))  
        beta<-beta+epsilon*gradient  
        #print(beta)  
        #print(dim(beta))  
      }  
    }  
    squared_loss[j,]<-loss  
  }  
}
```

```

    loss_iter<-1/N*(response-t(as.matrix(predictors))%*%beta)^2
    #print(dim(loss_iter))
    loss[i]<-loss_iter
  }
  else{
    gradient<-2/N*(t(as.matrix(predictors))%*%(response-as.matrix(predictors)%*%beta))
    #print(gradient)
    #print(dim(gradient))
    beta<-beta+epsilon*gradient
    #print(beta)
    #print(dim(beta))
    loss_iter<-1/N*(t(response-as.matrix(predictors)%*%beta)%*%(response-as.matrix(predictors)%*%beta))
    #print(dim(loss_iter))
    loss[i]<-loss_iter
  }
}
squared_loss[j,]<-loss
#print(loss)
}
#print(squared_loss)

return(squared_loss)
}

```

Here we generate the data to test the function

```

set.seed(2021)
sigma<-0.01
X<-mvrnorm(n=100, mu=rep(0,10),Sigma=diag(rep(1,10)))
a_star<-mvrnorm(n=1, mu=rep(0,10),Sigma=diag(rep(1,10)))
delta<-mvrnorm(n=1, mu=rep(0,100),Sigma=diag(rep(sigma^2,100)))

y<-X%*%a_star+delta


result_1<-stochastic_gd(X,y,0.01,1,10)
result_5<-stochastic_gd(X,y,0.01,5,10)
result_10<-stochastic_gd(X,y,0.01,10,10)

epochs<-1:10
loss_epochs_1<-result_1[,dim(result_1)[2]]
loss_epochs_5<-result_5[,dim(result_5)[2]]
loss_epochs_10<-result_10[,dim(result_10)[2]]

plot(x=epochs,y=loss_epochs_1,type="b",col="red",ylim=c(0,1),ylab="squared-error loss",main="Simulated")
points(x=epochs,y=loss_epochs_5,type="b",col="blue")
points(x=epochs,y=loss_epochs_10,type="b",col="green")
legend(8,1, legend=c("B=1", "B=5", "B=10"),
      col=c("red", "green", "blue"), lty=1, cex=0.8)

```

Simulated Data (sigma=0.01)



Here we change the value of sigma

```
set.seed(2021)
sigma<-1
X<-mvrnorm(n=100, mu=rep(0,10),Sigma=diag(rep(1,10)))
a_star<-mvrnorm(n=1, mu=rep(0,10),Sigma=diag(rep(1,10)))
delta<-mvrnorm(n=1, mu=rep(0,100),Sigma=diag(rep(sigma^2,100)))


y<-X%*%a_star+delta

result_1<-stochastic_gd(X,y,0.01,1,10)
result_5<-stochastic_gd(X,y,0.01,5,10)
result_10<-stochastic_gd(X,y,0.01,10,10)

epochs<-1:10
loss_epochs_1<-result_1[,dim(result_1)[2]]
loss_epochs_5<-result_5[,dim(result_5)[2]]
loss_epochs_10<-result_10[,dim(result_10)[2]]

plot(x=epochs,y=loss_epochs_1,type="b",col="red",ylim=c(0,1.5),ylab="squared-error loss", main="Simulated Data (sigma=0.01)",
points(x=epochs,y=loss_epochs_5,type="b",col="blue")
points(x=epochs,y=loss_epochs_10,type="b",col="green")
legend(8,1.2, legend=c("B=1", "B=5", "B=10"),
col=c("red", "green", "blue"), lty=1, cex=0.8)
```

Simulated Data (sigma=1)



Question 8